



Algorithm Efficiency and Sorting





How to Compare Different Problems and Solutions

- Two different problems
 - Which is harder/more complex?
- Two different solutions to the same problem
 - Which is better?
- Questions:
 - How can we compare different problems and solutions?
 - What does it mean to say that one problem or solution is more simpler or more complex than another?





Possible Solutions

- Idea: Code the solutions and compare them
 - Issues: machine, implementation, design, compiler, test cases, ...
- Better idea: Come up with a *machine- and implementation-independent* representation
 - # of steps
 - Time to do each step
- Use this representation to compare problems and solutions





Example: Traversing a Linked List

1. Node curr = head; // time: c_1
2. while(curr != null) { // time: c_2
3. System.out.println(curr.getItem());
4. curr=curr.getNext(); // time: c_3
5. }

- Given n elements in the list, total time =

$$1 \times c_1 + (n + 1) \times c_2 + n \times c_3$$

$$= n \times (c_2 + c_3) + c_2 + 1$$

$$= n \times d_1 + d_2$$

$$\propto n$$





Example: Nested Loops

```
1. for(i = 0; i < n; i++) {  
2.     for(j = 0; j < n; j++) {  
3.         System.out.println(i*j); // time: c  
4.     }  
5. }
```

■ Total time = $n \times n \times c$

$$\propto n^2$$





Example: Nested Loops II

1. `for(i = 0; i < n; i++) {`
2. `for(j = 0; j < i; j++) {`
3. `System.out.println(i*j);` // time: c
4. `}`
5. `}`

■ Total time = $\sum_{i=1}^n i \times c = c \sum_{i=1}^n i$

$$= c \times n \times (n-1) / 2$$
$$= d \times (n^2 - n)$$
$$\propto n^2 - n$$






Results

- Which algorithm is better?
 - Algorithm A takes $n^2 - 37$ time units
 - Algorithm B takes $n + 45$ time units
- Key Question: What happens as n gets large?
- Why?
 - Because for small n you can use any algorithm
 - Efficiency usually only matters for large n
- Answer: Algorithm B is better for large n
- Unless the constants are large enough
 - n^2
 - $n + 1000000000000000$



Graphically





Big O notation: $O(n)$

- An algorithm $g(n)$ is proportional to $f(n)$ if $g(n) = c_1 f(n) + c_2$
 - where $c_1 \neq 0$
- If an algorithm takes time proportional to $f(n)$, we say the algorithm is **order $f(n)$** , or **$O(f(n))$**
- Examples
 - $n+5$ is $O(n)$
 - $(n^2 + 3)/2$ is $O(n^2)$
 - $5n^2 + 2n/17$ is $O(n^2 + n)$





Exact Definition of $O(f(n))$

- An algorithm A is $O(f(n))$
 - IF there exists k and n_0
 - SUCH THAT A takes at most $k \times f(n)$ time units
 - To solve a problem of size $n \geq n_0$
-
- Examples:
 - $n/5 = O(n)$: $k = 5$, $n_0 = 1$
 - $3n^2 + 7 = O(n^2)$: $k = 4$, $n_0 = 3$
-
- In general, toss out constants and lower-order terms,
and $O(f(n)) + O(g(n)) = O(f(n) + g(n))$





Relationships between orders

- $O(1) < O(\log_2 n)$
- $O(\log_2 n) < O(n)$
- $O(n) < O(n \log_2 n)$
- $O(n \log_2 n) < O(n^2)$
- $O(n^2) < O(n^3)$
- $O(n^x) < O(x^n)$, for all x and n





Intuitive Understanding of Orders

- $O(1)$ – Constant function, independent of problem size
 - Example: Finding the first element of a list
- $O(\log_2 n)$ – Problem complexity increases slowly as the problem size increases.
 - Squaring the problem size only doubles the time.
 - Characteristic: Solve a problem by splitting into constant fractions of the problem (e.g., throw away $\frac{1}{2}$ at each step)
 - Example: Binary Search.
- $O(n)$ – Problem complexity increases linearly with the size of the problem
 - Example: counting the elements in a list.





Intuitive Understanding of Orders

- $O(n \log_2 n)$ – Problem complexity increases a little faster than n
 - Characteristic: Divide problem into subproblems that are solved the same way.
 - Example: mergesort
- $O(n^2)$ – Problem complexity increases fairly fast, but still manageable
 - Characteristic: Two nested loops of size n
 - Example: Introducing everyone to everyone else, in pairs
- $O(2^n)$ – Problem complexity increases very fast
 - Generally unmanageable for any meaningful n
 - Example: Find all subsets of a set of n elements





Search Algorithms

- Linear Search is $O(n)$
 - Look at each element in the list, in turn, to see if it is the one you are looking for
 - Average case $n/2$, worst case n
- Binary Search is $O(\log_2 n)$
 - Look at the middle element m . If $x < m$, repeat in the first half of the list, otherwise repeat in the second half
 - Throw away half of the list each time
 - Requires that the list be in sorted order
 - Sorting takes $O(n \log_2 n)$
- Which is more efficient?





Sorting





Selection Sort

- For each element i in the list
 - Find the smallest element j in the rest of the list
 - Swap i and j
- What is the efficiency of Selection sort?
- The for loop has n steps (1 per element of the list)
- Finding the smallest element is a linear search that takes $n/4$ steps on average (why?)
- The loops are nested: $n \times n/2$ on average: $O(n^2)$





Bubble sort

- Basic idea: run through the array, exchanging values that are out of order
 - May have to make multiple “passes” through the array
 - Eventually, we will have exchanged all out-of-order values, and the list will be sorted
 - Easy to code!
- Unlike selection sort, bubble sort doesn’t have an outer loop that runs once for each item in the array
- Bubble sort works well with either linked lists or arrays



Bubble sort: code

```
boolean done = false;
while(!done) {
    done = true;
    for (j = 0; j < length - 1; j++)
    {
        if (arr[j] > arr[j+1]) {
            temp = arr[j];
            arr[j] = arr[j+1];
            arr[j+1] = temp;
            done = false;
        }
    }
}
```

- Code is very short and simple
- Will it ever finish?
 - Keeps going as long as at least one swap was made
 - How do we know it'll eventually end?
- Guaranteed to finish: finite number of swaps possible
 - Small elements “bubble” up to the front of the array
 - Outer loop runs at most $nItems-1$ times
- Generally not a good sort
 - OK if a few items slightly out of order






Bubble sort: running time

- How long does bubble sort take to run?
 - Outer loop can execute a maximum of $nItems-1$ times
 - Inner loop can execute a maximum of $nItems-1$ times
- Answer: $O(n^2)$
 - Best case time could be much faster
 - Array nearly sorted would run very quickly with bubble sort
- Beginning to see a pattern: sorts seem to take time proportional to n^2
 - Is there any way to do better?
 - Let's check out insertion sort



What is insertion sort?



- Insertion sort: place the next element in the unsorted list where it “should” go in the sorted list
 - Other elements may need to shift to make room
 - May be best to do this with a linked list...





Pseudocode for insertion sort

```
while (unsorted list not empty) {  
  pop item off unsorted list  
  for (cur = sorted.first;  
    cur is not last && cur.value < item.value;  
    cur = cur.next) {  
    ;  
  }  
  if (cur.value < item.value) {  
    insert item after cur // last on list  
  } else {  
    insert item before cur  
  }  
}
```





How fast is insertion sort?

- Insertion sort has two nested loops
 - Outer loop runs once for each element in the original unsorted loop
 - Inner loop runs through sorted list to find the right insertion point
 - Average time: 1/2 of list length
- The timing is similar to selection sort: $O(n^2)$
- Can we improve this time?
 - Inner loop has to find element just past the one we want to insert
 - We know of a way to this in $O(\log n)$ time: binary search!
 - Requires arrays, but insertion sort works best on linked lists...
 - Maybe there's hope for faster sorting






How can we write faster sorting algorithms?

- Many common sorts consist of nested loops ($O(n^2)$)
 - Outer loop runs once per element to be sorted
 - Inner loop runs once per element that hasn't yet been sorted
 - Averages half of the set to be sorted
 - Examples
 - Selection sort
 - Insertion sort
 - Bubble sort
- Alternative: recursive sorting
 - Divide set to be sorted into two pieces
 - Sort each piece recursively
 - Examples
 - Mergesort
 - Quicksort



Sorting by merging: mergesort

1. Break the data into two equal halves
2. Sort the halves
3. Merge the two sorted lists
 - Merge takes $O(n)$ time
 - 1 compare and insert per item
 - How do we sort the halves?
 - Recursively
 - How many levels of splits do we have?
 - We have $O(\log n)$ levels!
 - Each level takes time $O(n)$
 - $O(n \log n)$!



Mergesort: the algorithm

```
void mergesort (int arr[], int sz) {  
    int half = sz/2;  
    int *arr2;  
    int k1, k2, j;  
    if (sz == 1) {  
        return;  
    }  
    arr2 = (int *)malloc(sizeof (int) * sz);  
    bcopy (arr, arr2, sz*sizeof(int));  
    mergesort (arr2, half);  
    mergesort (arr2+half, sz-half);  
    for (j=0, k1=0, k2=half; j < sz; j++) {  
        if ((k1 < half) && ((k2 >= sz) || (arr2[k1] < arr2[k2]))) {  
            arr[j] = arr2[k1++];  
        } else {  
            arr[j] = arr2[k2++];  
        }  
    }  
    free (arr2);  
}
```

Any array of size 1 is sorted!

Make a copy of the data to sort

Recursively sort each half

Merge the two halves

Use the item from first half if any left and


- There are no more in the second half or
- The first half item is smaller

Free the duplicate array



How well does mergesort work?

- Code runs in $O(n \log n)$
 - $O(n)$ for each “level”
 - $O(\log n)$ levels
- Depending on the constant, it may be faster to sort small arrays (1–10 elements or so) using an n^2 sort





Problems with mergesort

- Mergesort requires two arrays
 - Second array dynamically allocated (in C)
 - May be allocated on stack in C++
`int arr2[sz];`
 - This can take up too much space for large arrays!
- Mergesort is recursive
- These two things combined can be real trouble
 - Mergesort can have $\log n$ recursive calls
 - Each call requires $O(n)$ space to be allocated
- Can we eliminate this need for memory?






Solution: mergesort “in place”

- Mergesort builds up “runs” of correctly ordered items and then merges them
- Do this “in place” using linked lists
 - Eliminates extra allocation
 - Eliminates need for recursion (!)
- Keep two lists, each consisting of runs of 1 or more elements in sorted order
 - Combine the runs at the head of the lists into a single (larger) run
 - Place the run at the back of one of the lists
 - Repeat until you’re done



Mergesort “in place” in action



- Boxes with same color are in a single “run”
 - Specific color has no other meaning
- Runs get larger as the algorithm runs
 - Eventually, entire set is in one run!
- Algorithm works well with linked lists
 - No need to allocate extra arrays for merging!





Benefits of mergesort “in place”

- Algorithm may complete faster than standard mergesort
 - Requires fewer iterations if array is nearly sorted (lots of long runs)
 - Even small amounts of order make things faster
- No additional memory need be allocated
- No recursion!
 - Recursion can be messy if large arrays are involved
- Works well with linked lists
 - Standard mergesort is tougher with linked lists: need to find the “middle” element in a list
- May be less copying: simply rearrange lists






Quicksort: another recursive sort

- “Standard” mergesort requires too much memory
 - Extra array for merging
- Alternative: use quicksort
- Basic idea: partition array into two (possibly unequal) halves using a *pivot* element
 - Left half is all less than pivot
 - Right half is all greater than pivot
- Recursively continue to partition each half until array is sorted
 - Elements in a partition may move relative to one another during recursive calls
 - Elements can’t switch partitions during recursion



How quicksort works



- Pick a pivot element
- Divide the array to be sorted into two halves
 - Less than pivot
 - Greater than pivot
 - Need not be equal size!
- Recursively sort each half
 - Recursion ends when array is of size 1
 - Recursion may instead end when array is “small”: sort using traditional $O(n^2)$ sort
- How is pivot picked?
- What does algorithm look like?





Quicksort: pseudocode

```
quicksort (int theArray[], int nElem)
{
    if (nElem <= 1) // We're done
        return;
    Choose a pivot item p from theArray[]
    Partition the items of theArray about p
    Items less than p precede it
    Items greater than p follow it
    p is placed at index pIndex
    // Sort the items less than p
    quicksort (theArray, pIndex);
    // Sort the items greater than p
    quicksort (theArray+pIndex+1, nElem-(pIndex+1));
}
```

Key question: how do we pick a “good” pivot (and what makes a good pivot in the first place)?





Picking a pivot

- Ideally, a pivot should divide the array in half
 - How can we pick the middle element?
- Solution 1: look for a “good” value
 - Halfway between max and min?
 - This is slow, but can get a good value!
 - May be too slow...
- Solution 2: pick the first element in the array
 - Very fast!
 - Can result in slow behavior if we’re unlucky
- Most implementations use method 2





Quicksort: code

```
quicksort (int theArray[ ], int nElem)
{
    int pivotElem, cur, tmp;
    int endS1 = 0;
    if (nElem <= 1) return;
    pivotElem = theArray[0];
    for (cur = 1; cur < nElem; cur++) {
        if (theArray[cur] < pivotElem) {
            tmp = theArray[++endS1];
            theArray[endS1] = theArray[cur];
            theArray[cur] = tmp;
        }
    }
    theArray[0] = theArray[endS1];
    theArray[endS1] = pivotElem;
    quicksort (theArray, endS1); // Sort the two parts of the array
    quicksort (theArray+endS1+1, nElem-(endS1+1));
}
```





How fast is quicksort?

- Average case for quicksort: pivot splits array into (nearly) equal halves
 - If this is true, we need $O(\log n)$ “levels” as for mergesort
 - Total running time is then $O(n \log n)$
- What about the worst case?
 - Pick the minimum (or maximum) element for the pivot
 - S_1 (or S_2) is empty at each level
 - This reduces partition size by 1 at each level, requiring $n-1$ levels
 - Running time in the worst case is $O(n^2)$!
- For average case, quicksort is an excellent choice
 - Data arranged randomly when sort is called
 - May be able to ensure average case by picking the pivot intelligently
 - No extra array necessary!





Radix Sort: $O(n)$ (sort of)

- Equal length strings
- Group string according to last letter
- Merge groups in order of last letter
- Repeat with next-to-last letter, etc.
- Let's discuss how to do this
- Time: $O(nd)$
 - If d is constant (16-bit integers, for example), then radix sort takes $O(n)$

